



Introduction to CUDA 1 of 2

Patrick Cozzi
University of Pennsylvania
CIS 565 - Fall 2012

Announcements

- Project 0 due Tues 09/18
- Project 1 released today. Due Friday 09/28
 - Starts student blogs
 - Get approval for all third-party code

Acknowledgements

- Many slides are from David Kirk and Wen-mei Hwu's UIUC course:
 - <http://courses.engr.illinois.edu/ece498/al/>

GPU Architecture Review

- GPUs are specialized for
 - Compute-intensive, highly parallel computation
 - Graphics!
- Transistors are devoted to:
 - Processing
 - Not:
 - Data caching
 - Flow control

GPU Architecture Review

Transistor Usage



Figure 1-2. The GPU Devotes More Transistors to Data Processing

Image from: http://developer.download.nvidia.com/compute/cuda/3_2_prod/toolkit/docs/CUDA_C_Programming_Guide.pdf

Let's program this thing!

GPU Computing History

- 2001/2002 – researchers see GPU as data-parallel coprocessor
 - The *GPGPU* field is born
- 2007 – NVIDIA releases CUDA
 - *CUDA* – Compute Uniform Device Architecture
 - GPGPU shifts to *GPU Computing*
- 2008 – Khronos releases *OpenCL* specification

CUDA Abstractions

- A hierarchy of thread groups
- Shared memories
- Barrier synchronization

CUDA Terminology

- **Host** – typically the CPU
 - Code written in ANSI C
- **Device** – typically the GPU (data-parallel)
 - Code written in *extended* ANSI C
- Host and device have separate memories
- CUDA Program
 - Contains both host and device code

CUDA Terminology

- **Kernel** – data-parallel function
 - Invoking a kernel creates lightweight threads on the device
 - Threads are generated and scheduled with hardware
- Similar to a *shader* in OpenGL?

CUDA Kernels

- Executed N times in parallel by N different *CUDA threads*

```

// Kernel definition
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
}

int main()
{
    ...
    // Kernel invocation with N threads
    VecAdd<<<1, N>>>>(A, B, C);
}
    
```

Declaration Specifier points to `__global__`

Thread ID points to `threadIdx.x`

Execution Configuration points to `<<<1, N>>>>`

CUDA Program Execution




Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Thread Hierarchies

- **Grid** – one or more thread blocks
 - 1D or 2D
- **Block** – array of threads
 - 1D, 2D, or 3D
 - Each block in a grid has the same number of threads
 - Each thread in a block can
 - Synchronize
 - Access shared memory

Thread Hierarchies

- A **thread block** is a batch of threads that can **cooperate** with each other by:
 - Synchronizing their execution
 - For hazard-free shared memory accesses
 - Efficiently sharing data through a low latency **shared memory**
- Two threads from two different blocks cannot cooperate

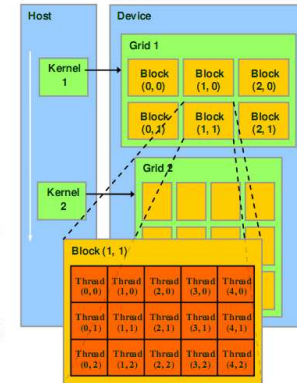




Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Thread Hierarchies

- **Block** – 1D, 2D, or 3D
 - Example: Index into vector, matrix, volume



Thread Hierarchies

- **Thread ID**: Scalar thread identifier
- **Thread Index**: `threadIdx`
- **1D**: Thread ID == Thread Index
- **2D** with size (D_x , D_y)
 - Thread ID of index (x , y) == $x + y D_y$
- **3D** with size (D_x , D_y , D_z)
 - Thread ID of index (x , y , z) == $x + y D_y + z D_x D_y$

Thread Hierarchies

```
// Kernel definition
__global__ void MatAdd(float A[N][N], float B[N][N],
                      float C[N][N])
{
    int i = threadIdx.x;
    int j = threadIdx.y;
    C[i][j] = A[i][j] + B[i][j];
}

int main()
{
    ...
    // Kernel invocation with one block of N * N * 1 threads
    int numBlocks = 1;
    dim3 threadsPerBlock(N, N);
    MatAdd<<numBlocks, threadsPerBlock>>>(A, B, C);
}
```

1 Thread Block

2D Index

2D Block

Thread Hierarchies

■ Thread Block

- Group of threads
 - G80 and GT200: Up to 512 threads
 - Fermi: Up to 1024 threads
- Reside on same processor core
- Share memory of that core

Thread Hierarchies

■ Thread Block

- Group of threads
 - G80 and GT200: Up to 512 threads
 - Fermi: Up to 1024 threads
- Reside on same processor core
- Share memory of that core




Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Thread Hierarchies

- Block Index: `blockIdx`
- Dimension: `blockDim`
 - 1D or 2D

Thread Hierarchies

```
// Kernel definition
__global__ void MatAdd(float A[N][N], float B[N][N],
                      float C[N][N])
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;
    if (i < N && j < N)
        C[i][j] = A[i][j] + B[i][j];
}

int main()
{
    ...
    // Kernel invocation
    dim3 threadsPerBlock(16, 16);
    dim3 numBlocks(N / threadsPerBlock.x, N / threadsPerBlock.y);
    MatAdd<<numBlocks, threadsPerBlock>>>(A, B, C);
}
```

16x16
Threads per block

2D Thread Block

Thread Hierarchies

■ Example: N = 32

- 16x16 threads per block (independent of N)
 - `threadIdx` ([0, 15], [0, 15])
- 2x2 thread blocks in grid
 - `blockIdx` ([0, 1], [0, 1])
 - `blockDim` = 16

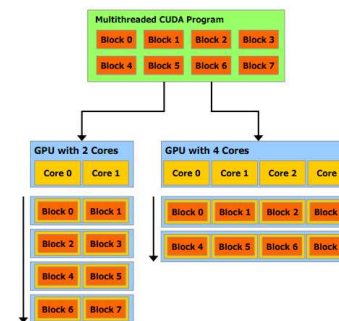
```
int i = blockIdx.x * blockDim.x + threadIdx.x;
int j = blockIdx.y * blockDim.y + threadIdx.y;
```

■ $i = [0, 1] * 16 + [0, 15]$

Thread Hierarchies

- Thread blocks execute independently
 - In any order: parallel or series
 - Scheduled in any order by any number of cores
 - Allows code to scale with core count

Thread Hierarchies



A multithreaded program is partitioned into blocks of threads that execute independently from each other, so that a GPU with more cores will automatically execute the program in less time than a GPU with fewer cores.

Figure 1-4. Automatic Scalability

Image from: http://developer.download.nvidia.com/compute/cuda/3_2_prod/toolkit/docs/CUDA_C_Programming_Guide.pdf

Thread Hierarchies

■ Threads in a block

- Share (limited) low-latency memory
- Synchronize execution
 - To coordinate memory accesses
 - `__syncThreads()`
 - Barrier – threads in block wait until all threads reach this
 - Lightweight




Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

• Device code can:

- R/W per-thread registers
- R/W per-thread local memory
- R/W per-block shared memory
- R/W per-grid global memory
- Read only per-grid constant memory

• Host code can

- R/W per grid global and constant memories

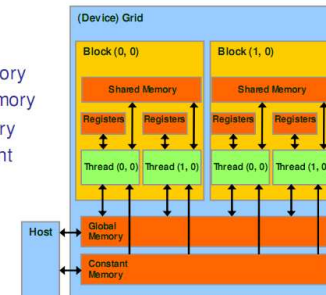


Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

■ Host can transfer to/from device

- *Global* memory
- *Constant* memory



Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

■ `cudaMalloc()`

- Allocate global memory on device

■ `cudaFree()`

- Frees memory




Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

```
float *Md;
int size = Width * Width * sizeof(float);

cudaMalloc((void**)&Md, size);
...
cudaFree(Md);
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

```
float *Md;
int size = Width * Width * sizeof(float);

cudaMalloc((void**)&Md, size);
...
cudaFree(Md);
```

Pointer to device memory

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

```
float *Md;
int size = Width * Width * sizeof(float);

cudaMalloc((void**)&Md, size);
...
cudaFree(Md);
```

Size in bytes

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

■ `cudaMemcpy()`

□ Memory transfer

- Host to host
- Host to device
- Device to host
- Device to device

Host

Device
Global Memory


- Similar to *buffer objects* in OpenGL

CUDA Memory Transfers

■ `cudaMemcpy()`

□ Memory transfer

- `Host to host`
- Host to device
- Device to host
- Device to device




CUDA Memory Transfers

■ `cudaMemcpy()`

□ Memory transfer

- Host to host
- `Host to device`
- Device to host
- Device to device




CUDA Memory Transfers

■ `cudaMemcpy()`

□ Memory transfer

- Host to host
- Host to device
- `Device to host`
- Device to device




CUDA Memory Transfers

■ `cudaMemcpy()`

□ Memory transfer


- Host to host
- Host to device
- Device to host
- `Device to device`



CUDA Memory Transfers

Host to device

```
cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);  
cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
```




Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

Destination (device) Source (host)


```
cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);  
cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
```



Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

CUDA Memory Transfers

```
cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);  
cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
```




Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply Reminder

- Vectors
- Dot products
- Row major or column major?
- Dot product per output element

Matrix Multiply



- $P = M * N$
- Assume M and N are square for simplicity


Image from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply

- 1,000 x 1,000 matrix
 - 1,000,000 dot products
 - Each 1,000 multiples and 1,000 adds

Matrix Multiply: CPU Implementation

```
void MatrixMulOnHost(float* M, float* N, float* P, int width)
{
    for (int i = 0; i < width; ++i)
        for (int j = 0; j < width; ++j)
        {
            float sum = 0;
            for (int k = 0; k < width; ++k)
            {
                float a = M[i * width + k];
                float b = N[k * width + j];
                sum += a * b;
            }
            P[i * width + j] = sum;
        }
}
```



Code from: <http://courses.engr.illinois.edu/ece498/al/lectures/lecture3%20cuda%20threads%20spring%202010.ppt>

Matrix Multiply: CUDA Skeleton

```
int main(void) {
    1. // Allocate and initialize the matrices M, N, P
    // I/O to read the input matrices M and N
    ....

    2. // M * N on the device
    MatrixMulOnDevice(M, N, P, width);

    3. // I/O to write the output matrix P
    // Free matrices M, N, P
    ...
    return 0;
}
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: CUDA Skeleton

```
int main(void) {
1. // Allocate and initialize the matrices M, N, P
  // I/O to read the input matrices M and N
  ....
2. // M * N on the device
  MatrixMulOnDevice(M, N, P, width);
  ....
3. // I/O to write the output matrix P
  // Free matrices M, N, P
  ....
  return 0;
}
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: CUDA Skeleton

```
int main(void) {
1. // Allocate and initialize the matrices M, N, P
  // I/O to read the input matrices M and N
  ....
2. // M * N on the device
  MatrixMulOnDevice(M, N, P, width);
  ....
3. // I/O to write the output matrix P
  // Free matrices M, N, P
  ....
  return 0;
}
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply

■ Step 1

- Add *CUDA memory transfers* to the skeleton

Matrix Multiply: Data Transfer

```
void MatrixMulOnDevice(float* M, float* N, float* P, int Width)
{
  int size = Width * Width * sizeof(float);

1. // Load M and N to device memory
  cudaMalloc(Md, size);
  cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
  cudaMalloc(Nd, size);
  cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

  // Allocate P on the device
  cudaMalloc(Pd, size);

2. // Kernel invocation code – to be shown later
  ....
3. // Read P from the device
  cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
  // Free device matrices
  cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}
```

← Allocate input

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: Data Transfer

```
void MatrixMulOnDevice(float* M, float* N, float* P, int Width)
{
    int size = Width * Width * sizeof(float);

    1. // Load M and N to device memory
    cudaMalloc(Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc(Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

    // Allocate P on the device
    cudaMalloc(Pd, size);

    2. // Kernel invocation code – to be shown later
    ...
    3. // Read P from the device
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
    // Free device matrices
    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}
```

Allocate output

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: Data Transfer

```
void MatrixMulOnDevice(float* M, float* N, float* P, int Width)
{
    int size = Width * Width * sizeof(float);

    1. // Load M and N to device memory
    cudaMalloc(Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc(Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

    // Allocate P on the device
    cudaMalloc(Pd, size);

    2. // Kernel invocation code – to be shown later
    ...
    3. // Read P from the device
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
    // Free device matrices
    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: Data Transfer

```
void MatrixMulOnDevice(float* M, float* N, float* P, int Width)
{
    int size = Width * Width * sizeof(float);

    1. // Load M and N to device memory
    cudaMalloc(Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc(Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

    // Allocate P on the device
    cudaMalloc(Pd, size);

    2. // Kernel invocation code – to be shown later
    ...
    3. // Read P from the device
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
    // Free device matrices
    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}
```

Read back from device

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: Data Transfer

```
void MatrixMulOnDevice(float* M, float* N, float* P, int Width)
{
    int size = Width * Width * sizeof(float);

    1. // Load M and N to device memory
    cudaMalloc(Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc(Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

    // Allocate P on the device
    cudaMalloc(Pd, size);

    2. // Kernel invocation code – to be shown later
    ...
    3. // Read P from the device
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
    // Free device matrices
    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}
```

■ Similar to GPGPU with GLSL.

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply

■ Step 2

- Implement the *kernel* in CUDA C

Matrix Multiply: CUDA Kernel

```
// Matrix multiplication kernel – thread specification
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int Width)
{
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue stores the Pd element that is computed by the thread
    float Pvalue = 0;

    for (int k = 0; k < Width; ++k)
    {
        float Mdelement = Md[ty * Md.width + k];
        float Ndelement = Nd[k * Nd.width + tx];
        Pvalue += Mdelement * Ndelement;
    }

    // Write the matrix to device memory each thread writes one element
    Pd[ty * Width + tx] = Pvalue;
}
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: CUDA Kernel

```
// Matrix multiplication kernel – thread specification
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int Width)
{
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue stores the Pd element that is computed by the thread
    float Pvalue = 0;

    for (int k = 0; k < Width; ++k)
    {
        float Mdelement = Md[ty * Md.width + k];
        float Ndelement = Nd[k * Nd.width + tx];
        Pvalue += Mdelement * Ndelement;
    }

    // Write the matrix to device memory each thread writes one element
    Pd[ty * Width + tx] = Pvalue;
}
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: CUDA Kernel

```
// Matrix multiplication kernel – thread specification
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int Width)
{
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue stores the Pd element that is computed by the thread
    float Pvalue = 0;

    for (int k = 0; k < Width; ++k)
    {
        float Mdelement = Md[ty * Md.width + k];
        float Ndelement = Nd[k * Nd.width + tx];
        Pvalue += Mdelement * Ndelement;
    }

    // Write the matrix to device memory each thread writes one element
    Pd[ty * Width + tx] = Pvalue;
}
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply: CUDA Kernel

```
// Matrix multiplication kernel – thread specification
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int Width)
{
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue stores the Pd element that is computed by the thread
    float Pvalue = 0;

    for (int k = 0; k < Width; ++k)
    {
        float Mdelement = Md[ty * Md.width + k];
        float Ndelement = Nd[k * Nd.width + tx];
        Pvalue += Mdelement * Ndelement;
    }

    // Write the matrix to device memory each thread writes one element
    Pd[ty * Width + tx] = Pvalue;
}
```

No locks or synchronization, why?

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply

Step 3

- Invoke the *kernel* in CUDA C

Matrix Multiply: Invoke Kernel

```
// Setup the execution configuration
dim3 dimBlock(WIDTH, WIDTH);
dim3 dimGrid(1, 1);
```


One block with width
by width threads

```
// Launch the device computation threads!
MatrixMulKernel<<<dimGrid, dimBlock>>>>(Md, Nd, Pd);
```

Code from: <http://courses.engr.illinois.edu/ece498/al/textbook/Chapter2-CudaProgrammingModel.pdf>

Matrix Multiply

- One Block of threads compute matrix Pd
 - Each thread computes one element of Pd
- Each thread
 - Loads a row of matrix Md
 - Loads a column of matrix Nd
 - Perform one multiply and addition for each pair of Md and Nd elements
 - Compute to off-chip memory access ratio close to 1:1 (not very high)
- Size of matrix limited by the number of threads allowed in a thread block



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2009
ECE 498AL Spring 2010, University of Illinois, Urbana-Champaign

Slide from: <http://courses.engr.illinois.edu/ece498/al/lectures/lecture2%20cuda%20spring%2009.ppt>

60



Matrix Multiply

- What is the major performance problem with our implementation?
- What is the major limitation?