

Programming GPUs for database applications

- outsourcing index search operations



Tim Kaldewey

Research Staff Member – Database Technologies

IBM Almaden Research Center

tkaldew@us.ibm.com



Quo Vadis ?



+ **ORACLE[®]** special projects



Why Search ?

Honestly, how many times a day do you visit

Google™

YAHOO!®

?



Quo Vadis ?



+ **ORACLE[®]** special projects

+






Quo Vadis ?



+ **ORACLE®** special projects



+

=





Agenda

- Introduction
 - GPU & DB (search) ?
- GPU search
 - A first implementation – binary search
 - Conventional search algorithms & GPUs – a mismatch
 - Back to the drawing board:
 - P-ary search
 - Experimental evaluation
 - Why it works
- Conclusions



Database Workloads


- Data-intensive
- Processor performance is not a problem
- Sifting through large quantities of data fast enough is





DB Performance – Where does Time Go

- CPU? I/O? Memory ? ¹
 - 10% indexed range selection




¹ A. Ailamaki, et al. DBMSs on a modern processor: Where does time go? VLDB'99




DB Performance – Where does Time Go

- CPU? I/O? Memory ? ¹
 - 10% indexed range selection




- It's getting worse ²

Relative Performance




¹ A. Ailamaki, et al. DBMSs on a modern processor: Where does time go? VLDB'99

² David Yen. Opening Doors to the MultiCore Era. MultiCore Expo 2006



DB Performance – “It's the memory stupid!” ³





³ R. Sites. It's the memory, stupid! MicroprocessorReport, 10(10),1996



DB Performance – “It's the memory stupid!” ³

- And worse:
 - Growth rates of main memory size have outstripped the growth rates of structured data in the enterprise ⁴
 - Multiple GB main memory DB ...



³ R. Sites. It's the memory, stupid! MicroprocessorReport, 10(10),1996


⁴ K. Schlegel. Emerging Technologies Will Drive Self-Service Business Intelligence. Garter Report 2/08



The (Memory) Wall 5



Relative Performance 2



² David Yen. Opening Doors to the MultiCore Era. MultiCore Expo 2006


⁵ W.A.Wulf et al. Hitting the memory wall: implications of the obvious. SIGARCH - Computer Architecture News'95



The (Memory) Wall 5



Relative Performance 2



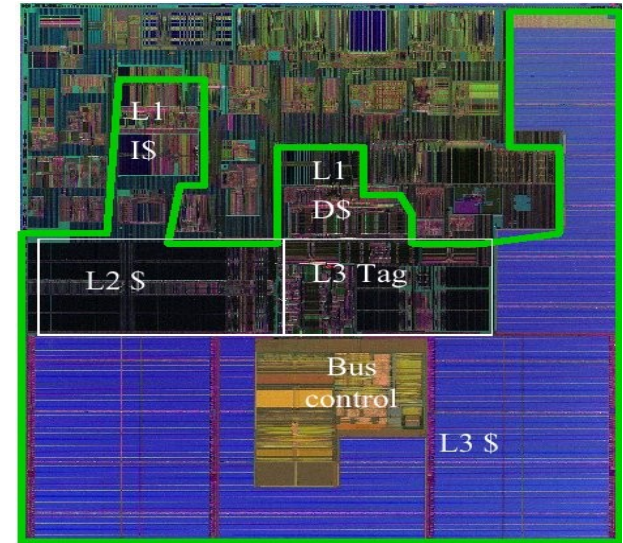
² David Yen. Opening Doors to the MultiCore Era. MultiCore Expo 2006

⁵ W.A.Wulf et al. Hitting the memory wall: implications of the obvious. SIGARCH - Computer Architecture News'95



Overcoming the Memory Wall

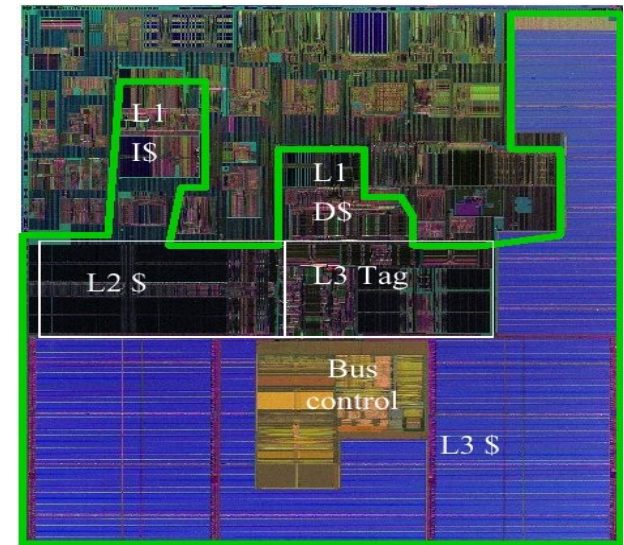
- Larger caches
 - Specialized processors
 - Top10 TPC-H – 6/10 use Itanium





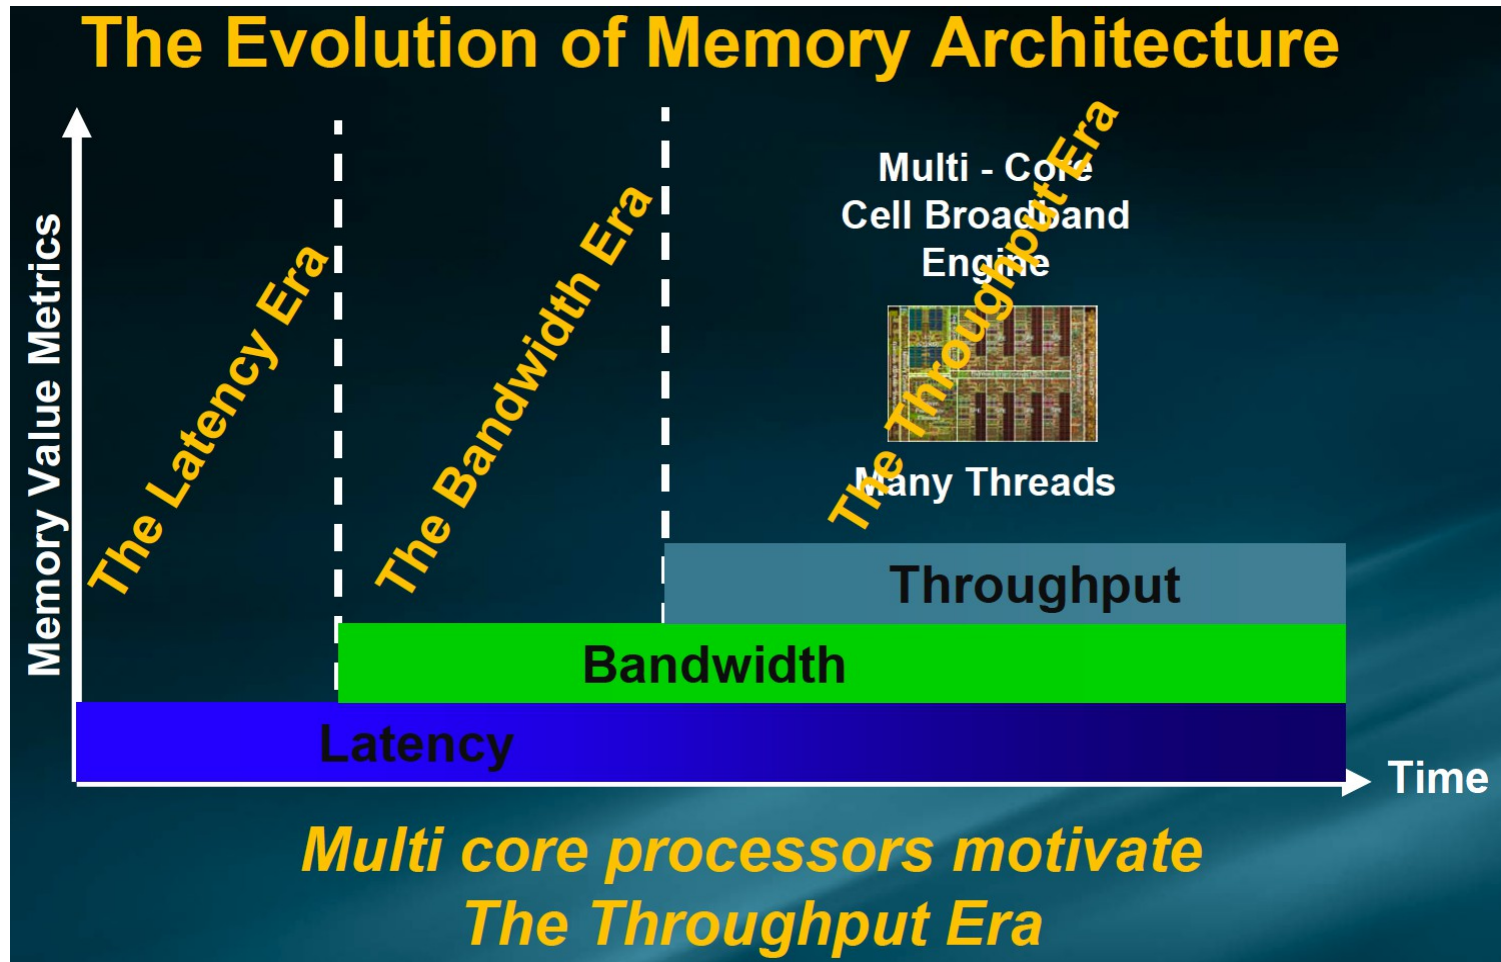
Overcoming the Memory Wall

- Larger caches
 - Specialized processors
 - Top10 TPC-H – 6/10 use Itanium
- Wait it out?





Parallel Memory Accesses → Throughput Computing



Source: Terabyte Bandwidth Initiative, Craig Hampel - Rambus, HotChips'08



GPUs as an example for highly parallel architectures

- Besides Teraflop(s) GPU's offer:
 - Massive Parallelism (240 cores)
 - 100+ GB/s memory bandwidth/throughput
 - Better performance per watt and per sqft. than CPUs






GPU performance specs & measurements



	GPU (GTX285)	CPU (i7-2600)
Power consumption	200 W	95 W
Peak Compute [Spec]	1063 GFLOP	109 GFLOP
Peak Memory Bandwidth [Spec]	160 GB/S	21 GB/s
Coalesced/Sequential Read [Measured]	140 GB/s	18 GB/s
Random Read [Measured]	8 GB/s	0.8 GB/s




GPU memory bandwidth – it's a **throughput** machine




Bandwidth of sequential (coalesced) 32-bit read access for multiple thread configurations. Results for a nVidia GTX 285 1.5GHz, GDDR3 1.2GHz.




GPU memory bandwidth




(a) coalesced (sequential) read



(b) random read



(c) coalesced (sequential) write



(d) random write

Parallel memory bandwidth for multiple thread configurations and access patterns. Results for a nVidia GTX 285 1.5GHz, GDDR3 1.2GHZ.



Agenda

- Introduction
 - GPU & DB (search) ?
- GPU search
 - A first implementation – binary search
 - Conventional search algorithms & GPUs – a mismatch
 - Back to the drawing board:
 - P-ary search
 - Experimental evaluation
 - Why it works
- Conclusions



A Simple implementation of (index) search

Keyword

Adam	1 , 2 , 3
Bethlehem	4 , 5
Character	1 , 2 , 3 , 301 , 5790
Drachenflieger	301 , 317 , 5790
Eva	1 , 2
Flughafenbahnhof	5790
Grabdenkmal	2 , 5790
Haubentaucher	300 , 5790



A Simple implementation of (index) search

Keyword

sorted
↓

Adam	1, 2, 3
Bethlehem	4, 5
Character	1, 2, 3, 301, 5790
Drachenflieger	301, 317, 5790
Eva	1, 2
Flughafenbahnhof	5790
Grabdenkmal	2, 5790
Haubentaucher	300, 5790

16 characters max.



A Simple implementation of (index) search

Keyword	
sorted ↓	Adam
	Bethlehem
	Character
	Drachenflieger
	Eva
	Flughafenbahnhof
	Grabdenkmal
	Haubentaucher
16 characters max.	

```
char indexCPU[4711];  
indexCPU[0]  
indexCPU[16]  
indexCPU[32]  
...
```



A Simple implementation of (index) search

Keyword	
sorted ↓	Adam
	Bethlehem
	Character
	Drachenflieger
	Eva
	Flughafenbahnhof
	Grabdenkmal
	Haubentaucher
16 characters max.	

```
char indexCPU[4711];  
indexCPU[0]  
indexCPU[16]  
indexCPU[32]  
...
```

- On the CPU we use a few library calls and we are done

```
char searchkey[16]= "Flughafenbahnhof";  
result = bsearch((void*)searchkey,indexCPU,  
                numentries,sizeof(char)*16,  
                (int (*)(const void*,const void*)) strcmp);
```



A Simple implementation of (index) search

Keyword	
sorted ↓	Adam
	Bethlehem
	Character
	Drachenflieger
	Eva
	Flughafenbahnhof
	Grabdenkmal
	Haubentaucher
16 characters max.	

```
char indexCPU[4711];
indexCPU[0]
indexCPU[16]
indexCPU[32]
...
```

- On the CPU we use a few library calls and we are done

```
char searchkey[16]= "Flughafenbahnhof";
result = bsearch((void*)searchkey,indexCPU,
                 numentries,sizeof(char)*16,
                 (int(*)(const void*,const void*)) strcmp);
```

- Can we just port a CPU implementation?



A Simple GPU implementation

- Get the data to the GPU

```
char* indexGPU;  
char* searchkeysGPU;  
char* resultsGPU;  
// copy the data  
cudaMalloc((void**)&indexGPU, sizeof(char)*wordlength*entries);  
cudaMemcpy(indexGPU, indexCPU, sizeof(char)*wordlength*entries,  
           CudaMemcpyHostToDevice);  
// copy the searchkey(s)  
cudaMalloc((void**)&searchkeysGPU, ...  
cudaMemcpy(searchkeysGPU, searchkeysCPU,  
           sizeof(char)*wordlength*numsearches,  
           CudaMemcpyHostToDevice);  
// make room for the results  
cudaMalloc((void**)&resultsGPU, ...
```



A Simple GPU implementation

- Get the data to the GPU

```
char* indexGPU;  
char* searchkeysGPU;  
char* resultsGPU;  
// copy the data  
cudaMalloc((void**)&indexGPU, sizeof(char)*wordlength*entries);  
cudaMemcpy(indexGPU, indexCPU, sizeof(char)*wordlength*entries,  
           CudaMemcpyHostToDevice);  
// copy the searchkey(s)  
cudaMalloc((void**)&searchkeysGPU, ...  
cudaMemcpy(searchkeysGPU, searchkeysCPU,  
           sizeof(char)*wordlength*numsearches,  
           CudaMemcpyHostToDevice);  
// make room for the results  
cudaMalloc((void**)&resultsGPU, ...
```

- Know your hardware (GTX 285, 32 SMs, 8 cores each, 240 cores)
 - Set up an execution configuration & call global function

```
dim3 Dg = dim3(30,0,0);  
dim3 Db = dim3(8,0,0);  
searchGPU<<< Dg,Db >>>(indexGPU, entries...
```




A Simple GPU implementation

- The GPU kernel

```
__global__ void searchGPU(char* index, int entries, int wordlength,
                          char* search_keys, int* results) {
    char* res;
    // use block and thread numbers for indexing
    res = bsearch(&search_keys[(blockIdx.x*BLOCK_SIZE)+threadIdx.x)
                  *wordlength],
                  index,
                  entries,
                  wordlength);
    // use block and thread numbers for indexing
    results[(blockIdx.x*BLOCK_SIZE)+threadIdx.x] = (res-data)/
                                                    MAX_WORD_LENGTH;
}
```



A Simple GPU implementation

- The GPU kernel

```
__global__ void searchGPU(char* index, int entries, int wordlength,
                          char* search_keys, int* results) {
    char* res;
    // use block and thread numbers for indexing
    res = bsearch(&search_keys[(blockIdx.x*BLOCK_SIZE)+threadIdx.x]
                  *wordlength],
                  index,
                  entries,
                  wordlength);
    // use block and thread numbers for indexing
    results[(blockIdx.x*BLOCK_SIZE)+threadIdx.x] = (res-data)/
                                                    MAX_WORD_LENGTH;
}
```

- There is no libc on the GPU =(
- Just stick `__device__` in front of the libc code?
- “bsearch” is recursive, but there is no recursion on the GPU
→ Write a iterative one ...



A Simple GPU binary search

```
__device__ char* bsearchGPU(char *key, char *base, int n, int size){
    char *mid_point;
    int cmp;

    while (n > 0) {
        mid_point = (char *)base + size * (n >> 1);
        if ((cmp = strcmpGPU(key, mid_point)) == 0)
            return (char *)mid_point;
        if (cmp > 0) {
            base = (char *)mid_point + size;
            n = (n - 1) >> 1;
        } // cmp < 0
        else n >>= 1;
    }
    return (char *)NULL;
}
```

- Still need strcmp



A Simple GPU binary search

```
__device__ char* bsearchGPU(char *key, char *base, int n, int size){
    char *mid_point;
    int cmp;

    while (n > 0) {
        mid_point = (char *)base + size * (n >> 1);
        if ((cmp = strcmpGPU(key, mid_point)) == 0)
            return (char *)mid_point;
        if (cmp > 0) {
            base = (char *)mid_point + size;
            n = (n - 1) >> 1;
        } // cmp < 0
        else n >>= 1;
    }
    return (char *)NULL;
}
```


- Still need strcmp
- Again, stick __device__ in front of the libc code

```
__device__ int strcmpGPU(char* s1, char* s2){
    while (*s1 == *s2++)
        if (*s1++ == 0) return 0;
    return (*s1 - *(s2 - 1));
}
```



Binary Search on the GPU

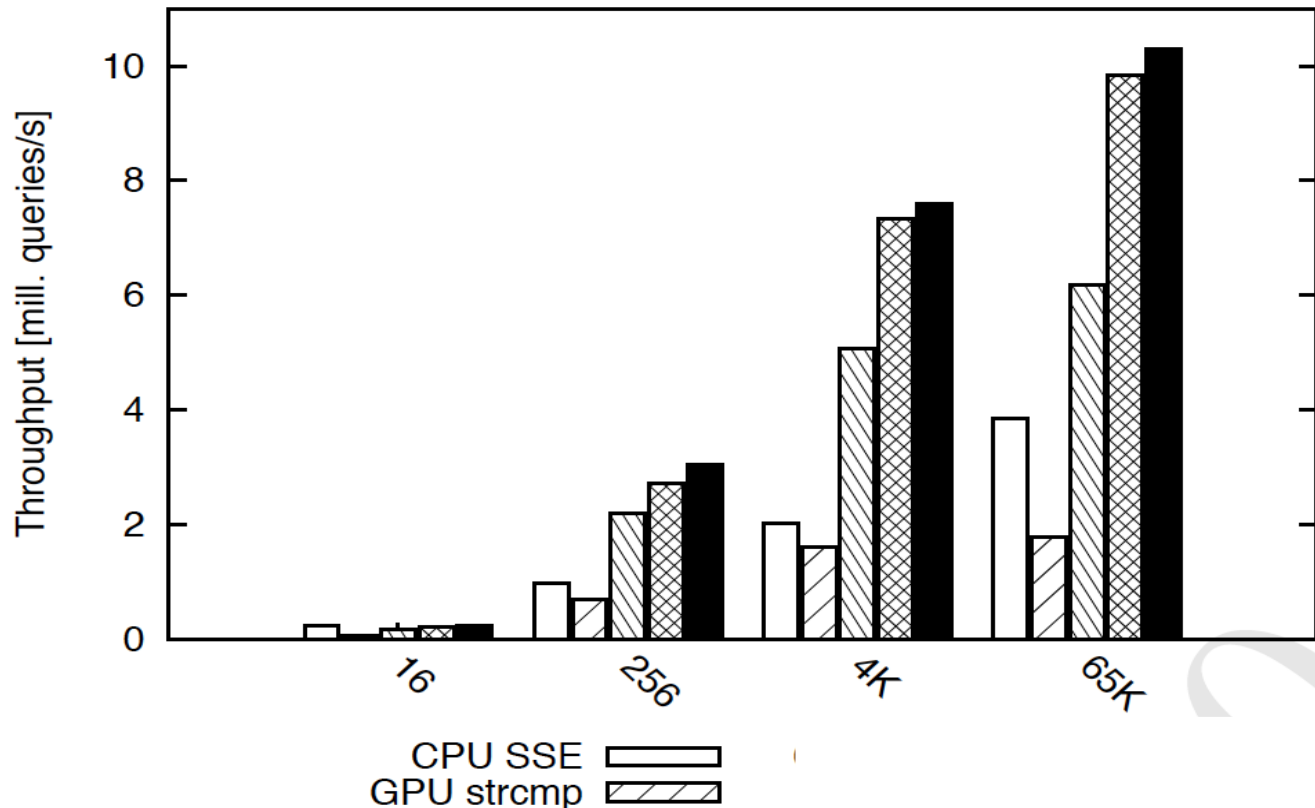
- Searching a large data set (512MB) with 33 million (225) 16-character strings





Binary Search on the GPU – Why is it slow?

- Searching a large data set (512MB) with 33 million (225) 16-character strings




- It's slower than a CPU implementation for all data set sizes!
 - Let's try some optimizations ...



Search requires to compare

- Search naturally requires MANY comparisons
- The strcmp() library function:

```
int strcmp(const char* s1, const char* s2){  
    while (*s1 == *s2++)  
        if (*s1++ == 0) return 0;  
    return (*s1 - *(s2 - 1));  
}
```






Search requires to compare

- Search naturally requires MANY comparisons
- The strcmp() library function:

```
int strcmp(const char* s1, const char* s2){  
    while (*s1 == *s2++)  
        if (*s1++ == 0) return 0;  
    return (*s1 - *(s2 - 1));  
}
```

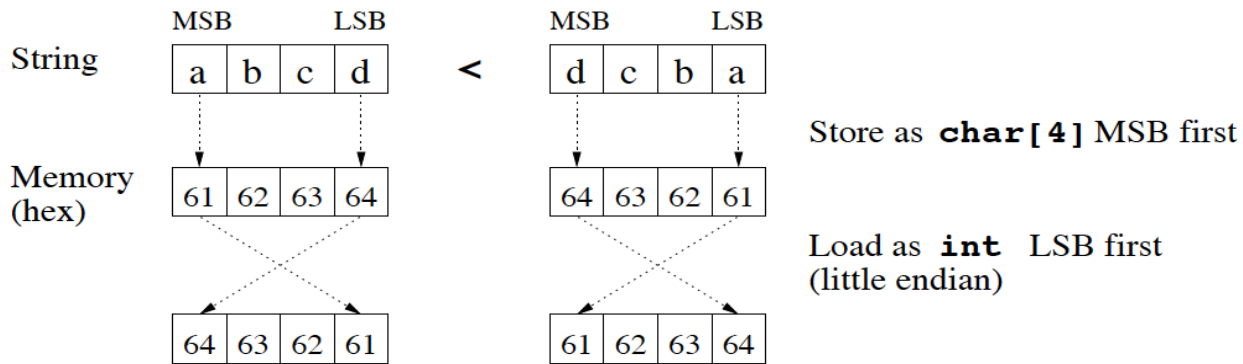


- **Byte-wise** memory access is known to be slow



Optimizing compare operations

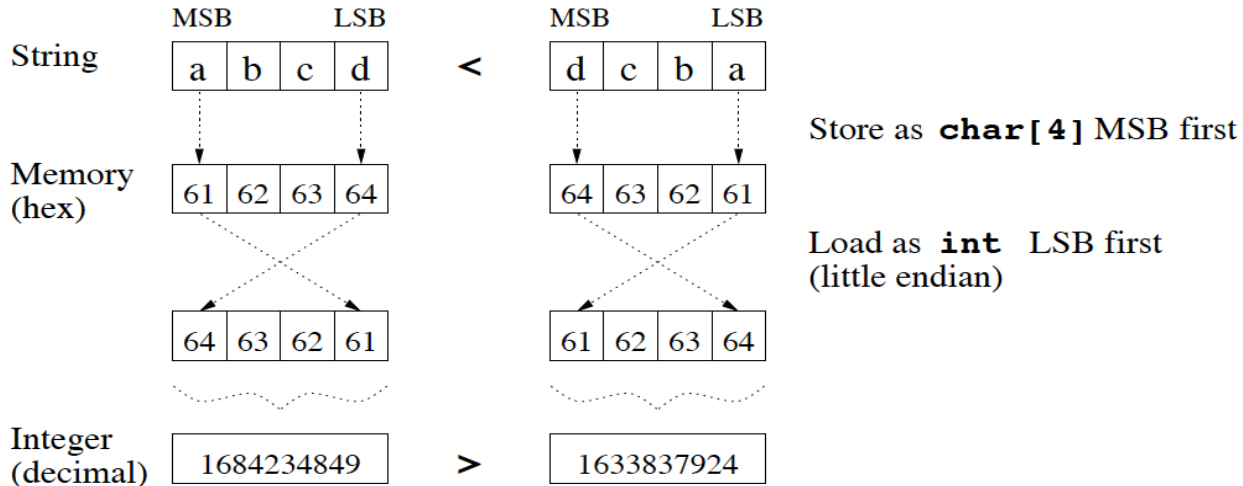
- How about vector string comparison, a la SSE?
- No Byte vectors on the GPU ... but Integer vectors





Optimizing compare operations

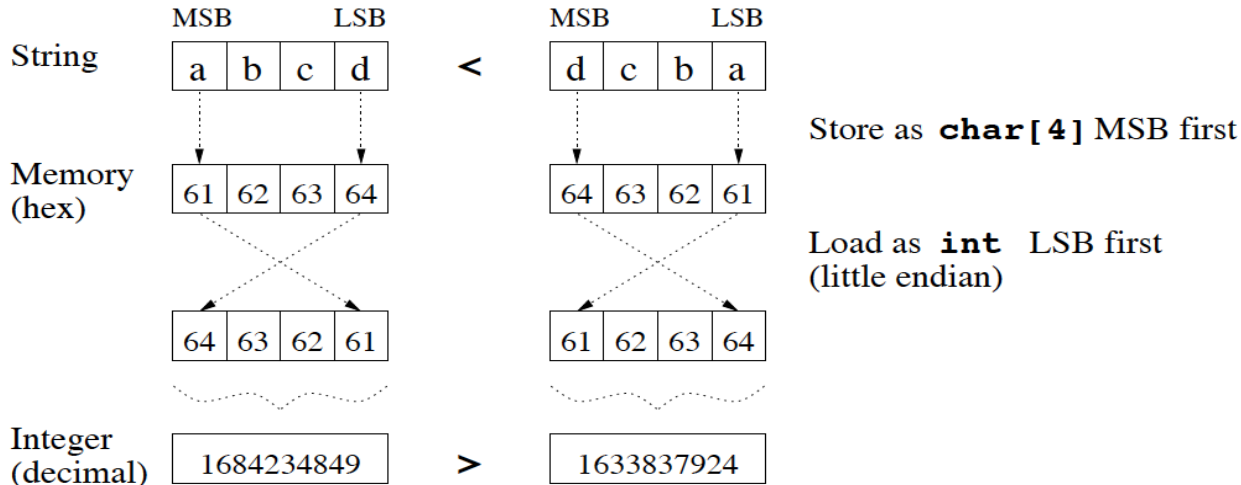
- How about vector string comparison, a la SSE?
- No Byte vectors on the GPU ... but Integer vectors





Optimizing compare operations

- How about vector string comparison, a la SSE?
- No Byte vectors on the GPU ... but Integer vectors



- Loading character strings as int changes endianness
- CPU has `bswap`, on the GPU we have to write it:

```
#define BSWP( x ) ; \
temp = ( x ) << 24 ; \
temp = temp | ( ( ( x ) << 8) & 0x00FF0000 ) ; \
temp = temp | ( ( ( unsigned ) ( x ) >> 8) & 0x0000FF00 ) ; \
x = temp | ( ( unsigned ) ( x ) >> 24 ) ;
```



Optimizing compare operations

- Comparing integer vectors (bswap for <> skipped for clarity)

```
__device__ int intcmp(uint4* a, uint4* b){


    int r =1;
    if ((*a).x < (*b).x)
        r=-1;
    else if ((*a).x == (*b).x) {
        if ((*a).y < (*b).y)
            r=-1;
        else if ((*a).y == (*b).y) {
            if ((*a).z < (*b).z)
                r=-1;
            else if ((*a).z == (*b).z) {
                if ((*a).w < (*b).w)
                    r=-1;
                else if ((*a).w == (*b).w)
                    r=0;
            }
        }
    }
    return r;
}
```

- Still dereferencing 16 memory pointers ...



Binary Search on the GPU – Why is it slow?

- Searching a large data set (512MB) with 33 million (225) 16-character strings



- With intcmp it's only marginally faster than a CPU implementation
- We still do pointer chasing, i.e. roundtrips to memory ...



Reducing global memory access

- Intcmp is memory latency sensitive

Processor	L1 [cyc]	L2 [cyc]	L3 [cyc]	mem [cyc]
Intel Core i7 2.6GHz	4	10	40	350
nVidia GT200b 1.5 GHz	4	n/a	n/a	500

- We can use shared memory like L1

x 16 for each
comparison !!!



Reducing global memory access

- Intcmp is memory latency sensitive

Processor	L1 [cyc]	L2 [cyc]	L3 [cyc]	mem [cyc]
Intel Core i7 2.6GHz	4	10	40	350
nVidia GT200b 1.5 GHz	4	n/a	n/a	500

- We can use shared memory like L1

x 16 for each
comparison !!!

```
__shared__ uint4 cache[NUM_THREADS*2];


__device__ uint4* bsearchGPU( uint4 *key,  uint4 *base,
                              size_t nmemb,  size_t size)
{
    uint4 *mid_point;
    int  cmp;
    cache[threadIdx.x*2]= *key;

    while (nmemb > 0) {
        mid_point = (uint4 *)base + size * (nmemb >> 1);
        cache[threadIdx.x*2+1]= *mid_point;
        if ((cmp = intcmp(&cache[threadIdx.x*2],
                        &cache[threadIdx.x*2+1]))== 0)
            return (uint4 *)mid_point;
    }
}
```



Binary Search on the GPU – optimized


- Searching a large data set (512MB) with 33 million (225) 16-character strings





GPU architecture reminder – SIMD/SIMT

- Inside Streaming Multiprocessor
 - Single Instruction Multiple Threads/Data (SIMT/SIMD)
 - All cores in 1SM execute same instruction or no-op (SIMD threads)
 - Warps of 32 threads (or more, to hide memory latency)





Multi-threaded Binary Search – Example

- 1 Index: a sorted char array 32 entries
- 4 queries: **t**, **8**, **f**, **r**
- 4 processor cores: P1–P4
- 1 processor core – 1 search: P0:**t**, P1:**8**, P2:**f**, P3:**r**
- Theoretical worst-case execution time: $\log_2(32)=5$

4	5	6	7	8	9	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
---	---	---	---	----------	---	---	---	---	---	---	----------	---	---	---	---	---	---	---	---	---	---	---	----------	---	----------	---	---	---	---	---	---

Multi-threaded Binary Search – Example

- 1 Index: a sorted char array 32 entries
- 4 queries: t , 8 , f , r
- 4 processor cores: P1–P4
- 1 processor core – 1 search: P0:t , P1:8 , P2:f , P3:r
- Theoretical worst-case execution time: $\log_2(32)=5$

Iter. 1)

4	5	6	7	8	9	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

P0:t, P1:8, P2:f, P3:r




Iter. 2)

4	5	6	7	8	9	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Diagram illustrating the state of the array after Iteration 2. The array contains 32 elements. The first 16 elements are grouped by a bracket labeled "P1:8, P2:f". The last 16 elements are grouped by a bracket labeled "P0:t, P3:r". The elements 8, f, r, and t are highlighted in red.



Multi-threaded Binary Search – Example





Conventional multi-threading – Analysis

- 100% utilization requires
#cores concurrent queries
- Queries finishing early
→ utilization < 100%
- Memory access collisions
→ serialized memory access
- #memory accesses $\log_2(n)$
- More threads
→ more results
→ response time likely to be
worst case: $\log_2(n)$



Can we improve the worst case?



Agenda

- Introduction
 - GPU & DB (search) ?
- Porting search to the GPU using CUDA
 - Conventional search and GPU architecture – a mismatch
 - Back to the drawing board:
 - P-ary search – the algorithm
 - Experimental evaluation
 - Why it works
- Conclusions



Our Goal

- Improve response time (latency) of core database functions like search in the era of throughput oriented (parallel) computing.



Research Question

- How can we (algorithmically) exploit parallelism to improve response time (of search)?
 - Can we trade-off throughput for latency?
 - Do we have to trade?

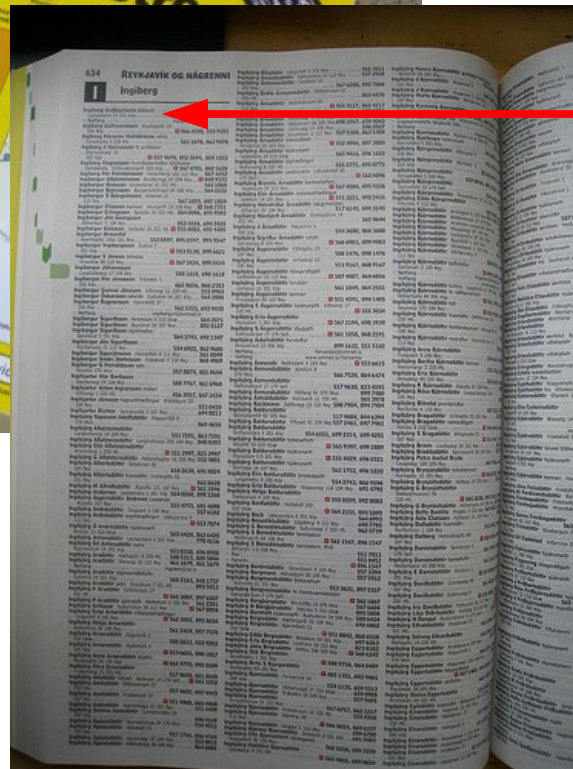


Binary Search

- How Do you (efficiently) search an index?



- Open phone book ~middle




- 1st name = whom you are looking for?
- $<, > ?$
- Iterate
 - Each iteration: $\#entries/2$ ($n/2$)
 - Total time: $\rightarrow \log_2(n)$



Parallel (Binary) Search

- What if you have some friends (3) to help you ?




- Divide et impera !
 - Each is using binary search takes $\log_2(n/4)$
 - All can work in parallel \rightarrow faster: $\log_2(n/4) < \log_2(n)$
- Give each of them $\frac{1}{4}$ *

* You probably want to tear it a little more intelligent than that, e.g. at the binding ;-)



Parallel (Binary) Search

- What if you have some friends (3) to help you ?





- Divide et impera !
 - Each is using binary search takes $\log_2(n/4)$
 - All can work in parallel \rightarrow faster: $\log_2(n/4) < \log_2(n)$
 - 3 of you are **wasting time** !
- Give each of them $\frac{1}{4}$ *

* You probably want to tear it a little more intelligent than that, e.g. at the binding ;-)



P-ary Search

- Divide et impera !!



...

- How do we know who has the right piece ?



P-ary Search

- Divide et impera !!



- How do we know who has the right piece ?



- It's a sorted list:
 - Look at first and last entry of a subset
 - If **first entry** < searched name < **last entry**
 - Redistribute
 - Otherwise ... throw it away
 - Iterate



P-ary Search

- What do we get?



+

- Each iteration: $n/4$
→ $\log_4(n)$
- Assuming redistribution time is negligible:
 $\log_4(n) < \log_2(n/4) < \log_2(n)$
- But each does 2 lookups !
- How time consuming are **lookup** and **redistribution** ?



P-ary Search

- What do we get?



+

- Each iteration: $n/4$
→ $\log_4(n)$
- Assuming redistribution time is negligible:
 $\log_4(n) < \log_2(n/4) < \log_2(n)$
- But each does 2 lookups !
- How time consuming are **lookup** and **redistribution** ?

||

memory
access

||

synchronization



P-ary Search

- What do we get



+

- Each iteration: $n/4$
→ $\log_4(n)$
- Assuming redistribution time is negligible:
 $\log_4(n) < \log_2(n/4) < \log_2(n)$
- But each does 2 lookups !
- How time consuming are lookup and redistribution ?

||

memory
access

||

synchronization


- Searching a database index can be implemented the same way
 - Friends = Processor cores (threads)
 - Without destroying anything ;-)



P-ary Search - Implementation

- Strongly relies on fast synchronization
 - friends = threads / processor cores / vector elements

Iteration 1)







P-ary Search - Implementation

- Strongly relies on fast synchronization
 - friends = threads / processor cores / vector elements

Iteration 1)




Iteration 2)





P-ary Search - Implementation

- Strongly relies on fast synchronization
 - friends = threads / processor cores / vector elements

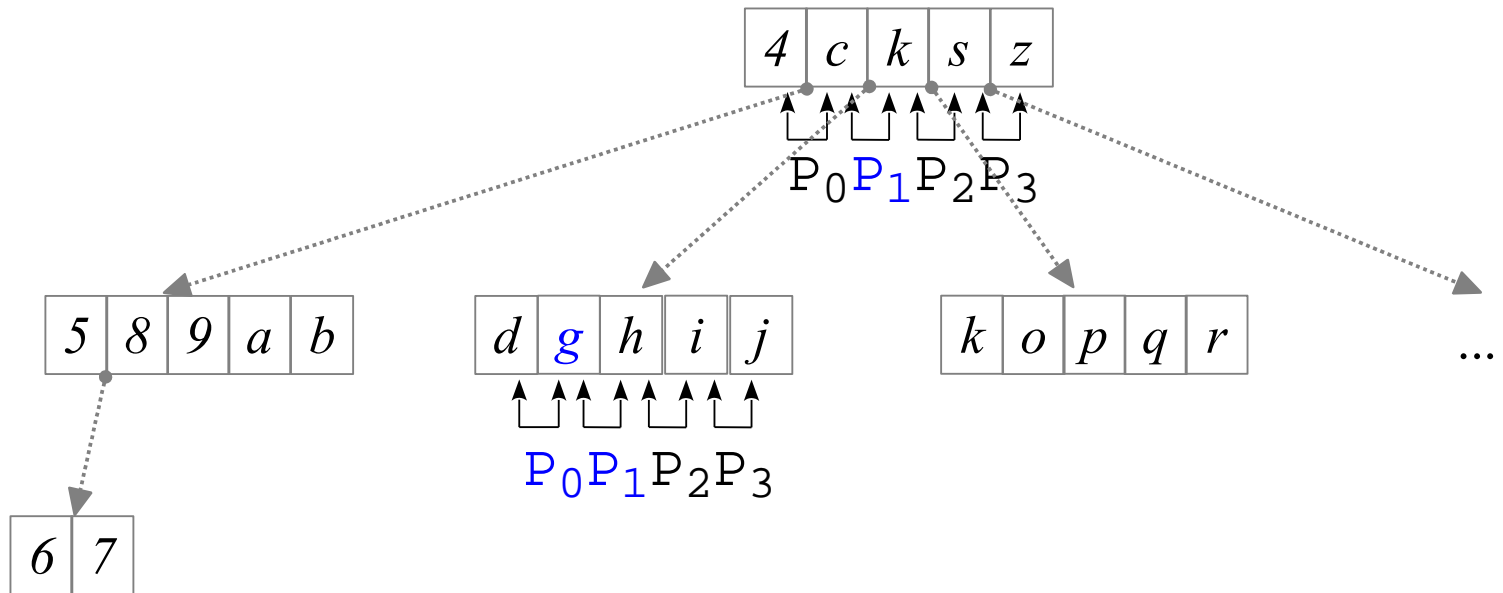


- Synchronization ~ repartition cost
pthreads (\$\$), `cmpxchg`(\$),
SIMD {SSE-vector, GPU threads via shared memory} (~0)
- Implementation using a B-tree is similar and (obviously) faster



P-ary Search - Implementation

- B-trees group pivot elements into nodes



- Access to pivot elements is coalesced instead of a gather
- Nodes can also be mapped to
 - Cache Lines (CSB+ trees)
 - Vectors (SSE)



P-ary Search on a sorted list – Implementation (1)

```
__global__ void parySearchGPU(int* data , int range_length , int*  
                             search keys , int* results)  
  
    int sk , old_range_length=range_length, range start ;  
    // initialize search range starting with the whole data set  
    // this is done by one thread  
    if (threadIdx.x==0) {  
        range_offset=0;  
        // cache search key and upper bound in shared memory  
        cache[BLOCKSIZE]=0x7FFFFFFF;  
        cache[BLOCKSIZE+1]=searchkeys[blockIdx.x];  
    }  
    // require a sync, since each thread is going to  
    // read the above now  
    Syncthreads();  
    sk = cache[BLOCKSIZE+1];
```



P-ary Search on a sorted list – Implementation (2)

```
// repeat until the #keys in the search range is
// smaller than the number of threads
while (range_length>BLOCKSIZE){
    // calculate search range for this thread
    // avoiding floating point operations
    range_length = range_length/BLOCKSIZE;
    if (range_length * BLOCKSIZE < old_range_length)
        range_length+=1;
    old_range_length=range_length;
    range_start = range_offset + threadIdx.x * range_length;

    // cache the boundary keys
    cache[threadIdx.x]=data[range_start];
    __syncthreads();
    // if the seached key is within this thread's subset,
    // make it the one for the next iteration
    if (sk>=cache[threadIdx.x] && sk<cache[threadIdx.x+1]){
        range_offset = range_start;
    }
    // all threads need to start next iteration
    // with the new subset
    __syncthreads();
}
```




P-ary Search on a sorted list – Implementation (3)

```
// last iteration
range_start = range_offset + threadIdx.x;
if (sk==data[range_start])
    results[blockIdx.x]=range_start;
}
```



P-ary Search – Analysis



- 100% processor utilization for each query
- Multiple threads can find a result
 - Does not change correctness





P-ary Search – Analysis



- 100% processor utilization for each query
- Multiple threads can find a result
 - Does not change correctness
- Convergence depends on #threads
GTX285: 1 SM, 8 cores(threads) $\rightarrow p=8$
- Better Response time
 - $\log_p(n)$ vs $\log_2(n)$





P-ary Search – Analysis



- 100% processor utilization for each query
- Multiple threads can find a result
 - Does not change correctness
- Convergence depends on #threads
GTX285: 1 SM, 8 cores(threads) $\rightarrow p=8$
- Better Response time
 - $\log_p(n)$ vs $\log_2(n)$
- More memory access
 - $(p*2 \text{ per iteration}) * \log_p(n)$
 - Caching
 - $(p-1) * \log_p(n)$ vs. $\log_2(n)$





P-ary Search – Analysis


- 100% processor utilization for each query
- Multiple threads can find a result
 - Does not change correctness
- Convergence depends on #threads
GTX285: 1 SM, 8 cores(threads) $\rightarrow p=8$
- Better Response time
 - $\log_p(n)$ vs $\log_2(n)$
- More memory access
 - $p \cdot 2$ per iteration $\cdot \log_p(n)$
 - Caching
 $(p-1) \cdot \log_p(n)$ vs. $\log_2(n)$
- Lower Throughput
 - $1/\log_p(n)$ vs $p/\log_2(n)$





P-ary Search (GPU) – Throughput

- Superior throughput compared to conventional algorithms




Searching a 512MB data set with 134mill. 4-byte integer entries,
Results for a nVidia GT200b, 1.5GHz, GDDR3 1.2GHz.



P-ary Search (GPU) – Response Time

- Response time is workload independent for B-tree implementation




Searching a 512MB data set with 134mill. 4-byte integer entries,
Results for a nVidia GT200b, 1.5GHz, GDDR3 1.2GHz.



P-ary Search (GPU) – Scalability

- GPU Implementation using SIMT (SIMD threads)
- Scalability with increasing #threads (P)




64K search queries against a 512MB data set with 134mill. 4-byte integer entries,
Results for a nVidia GT200b, 1.5GHz, GDDR3 1.2GHz.



P-ary Search (GPU) – Scalability

- GPU Implementation using SIMT (SIMD threads)
- Scalability with increasing #threads (P)




64K search queries against a 512MB data set with 134mill. 4-byte integer entries,
Results for a nVidia GT200b, 1.5GHz, GDDR3 1.2GHz.



P-ary Search(CPU) = K-ary Search

- K-ary¹ search is the same algorithm ported to the CPU using SSE vectors (int4) → convergence rate $\log_4(n)$




Searching a 512MB data set with 134mill. 4-byte integer entries,
Core i7 2.66GHz, DDR3 1666.



P-ary Search(CPU) = K-ary Search

- Throughput scales proportional to #threads



64K search queries against a 512MB data set with 134mill. 4-byte integer entries, Core i7 2.66GHz, DDR3 1666.



P-ary search - an architecture perspective

- Architecture trends
 - Memory latency has bottomed out more than a decade ago
 - Parallel memory bandwidth keeps increasing
 - e.g. Core 2 8GB/s, Core i7 24GB/s (10GB/s per core)
 - Multi-core is just the beginning, many-core is the future
 - Cache per core keeps decreasing (GPU, no caches)
 - Linear (coalesced) memory accesses take its place
 - Core/ thread synchronization costs keep decreasing
- ➔ Only thing to hope for are increases in **parallel** memory **bandwidth**



P-ary search - an architecture perspective

- Architecture trends
 - Memory latency has bottomed out more than a decade ago
 - Parallel memory bandwidth keeps increasing
 - e.g. Core 2 8GB/s, Core i7 24GB/s (10GB/s per core)
 - Multi-core is just the beginning, many-core is the future
 - Cache per core keeps decreasing (GPU, no caches)
 - Linear (coalesced) memory accesses take its place
 - Core/ thread synchronization costs keep decreasing
- ➔ Only thing to hope for are increases in **parallel** memory **bandwidth**
- P-ary search was designed under this premises and provides
 - Scalable performance – fast thread synchronization
 - Reduced query response time – parallel memory access
 - Increased throughput – coalesced memory access
 - Workload independent constant query execution time



Questions